



ToolCall M30 vs M60 comparison

470M-token standard baseline versus resumed SwiGLU + GQA run


Final validation loss




Validation learning curves



Training-loss convergence



Training runtime



Run	Status	Tokens seen	General loss	Structured loss	Train EMA	Time	Median speed	Peak VRAM
M30 standard	completed	469.0M	2.8507	2.0981	2.6629	3.21 h	40.4k tok/s	3.15 GiB
M60 SwiGLU + GQA (resumed)	paused	468.7M	2.3521	2.1073	2.5078	6.59 h	19.9k tok/s	2.79 GiB